

A diagram placeholder for HDFS Architecture. It consists of a light gray rectangular box in the center, flanked by two vertical blue lines. At the top of each blue line, there is a horizontal segment that extends outwards, forming a bracket-like shape.


HDFS ARCHITECTURE

HADOOP

- ✓ Apache Hadoop is a framework that allows for the distributed processing of large data sets across clusters of commodity computers using a simple programming model.

- ✓ Companies using Hadoop:

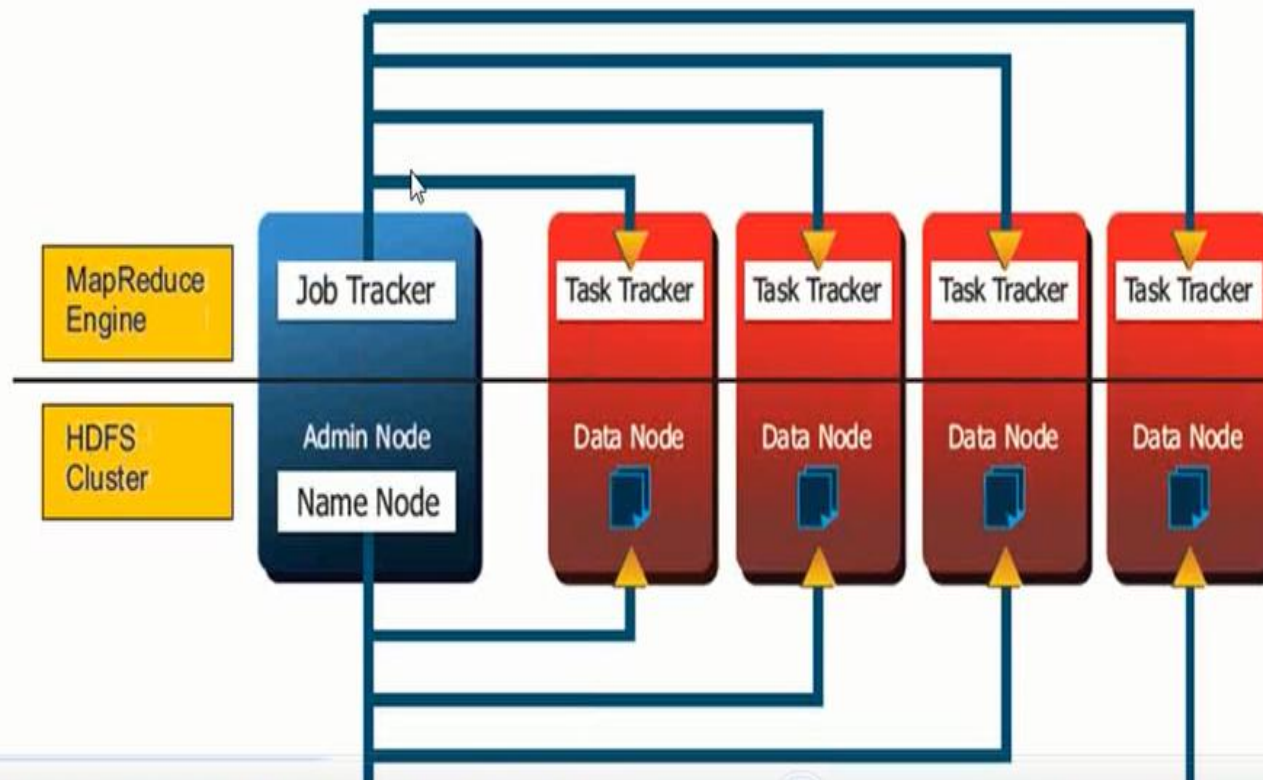
- Yahoo
- Google
- Facebook
- Amazon
- AOL
- IBM
- And many more at



<http://wiki.apache.org/hadoop/PoweredBy>

HDFS CORE COMPONENTS

- ✓ HDFS – Hadoop Distributed File System (storage)
- ✓ MapReduce (processing)



HDFS

HDFS - Hadoop Distributed File System

- ✓ Highly fault-tolerant
- ✓ High throughput
- ✓ Suitable for applications with large data sets
- ✓ Streaming access to file system data
- ✓ Can be built out of commodity hardware

DESIGN OF HDFS

HDFS is a file system designed for storing very large files with streaming data access patterns, running clusters on commodity hardware.

AREAS WHERE HDFS IS NOT GOOD

- Low-latency data access
- Lots of small files
- Multiple writers, arbitrary file modifications

HDFS COMPONENTS

- Namenodes
- Datanodes

MAIN COMPONENTS OF HDFS

✓ NameNode:

- ✓ master of the system
- ✓ maintains and manages the blocks which are present on the DataNodes




✓ DataNodes:

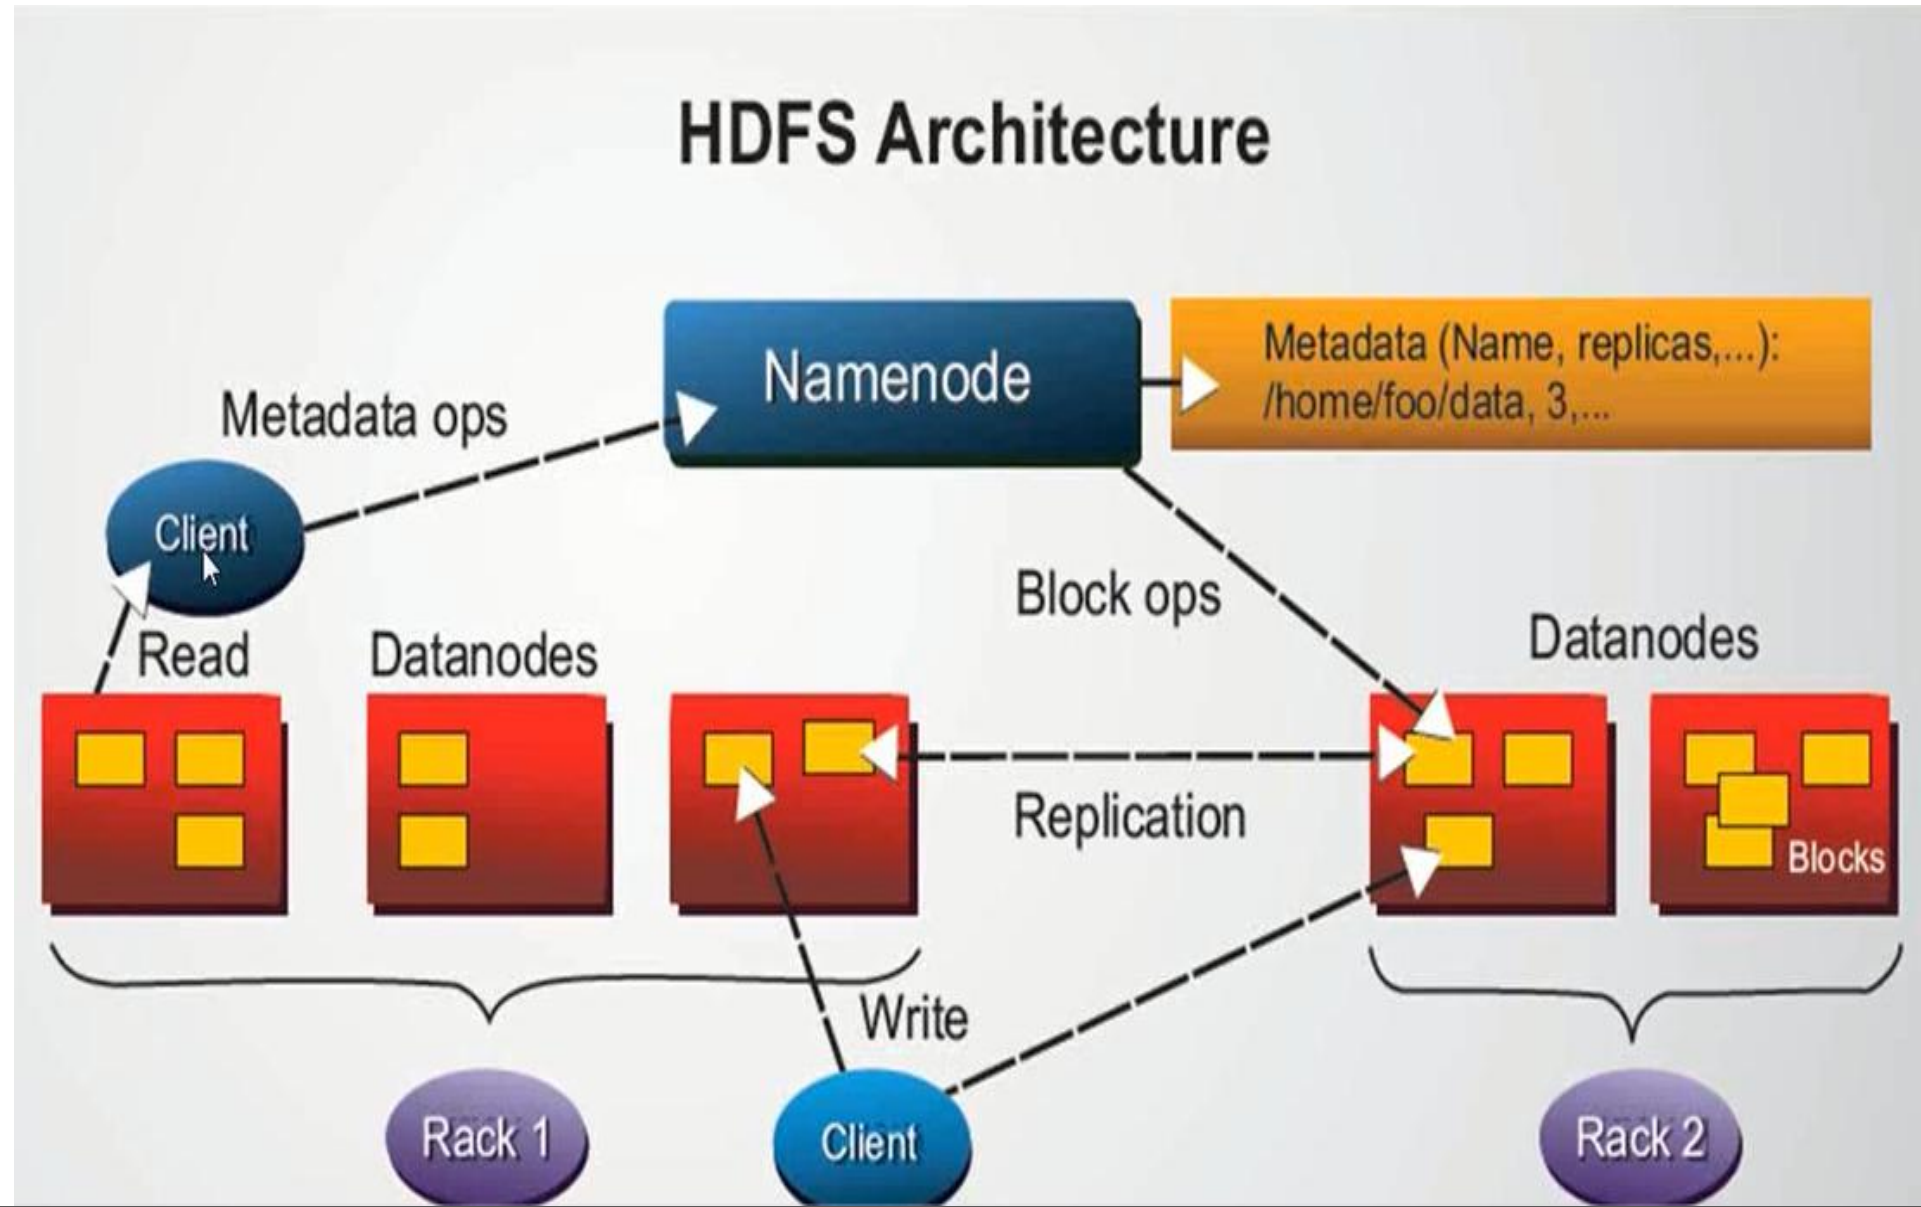
- ✓ slaves which are deployed on each machine and provide the actual storage
- ✓ responsible for serving read and write requests for the clients



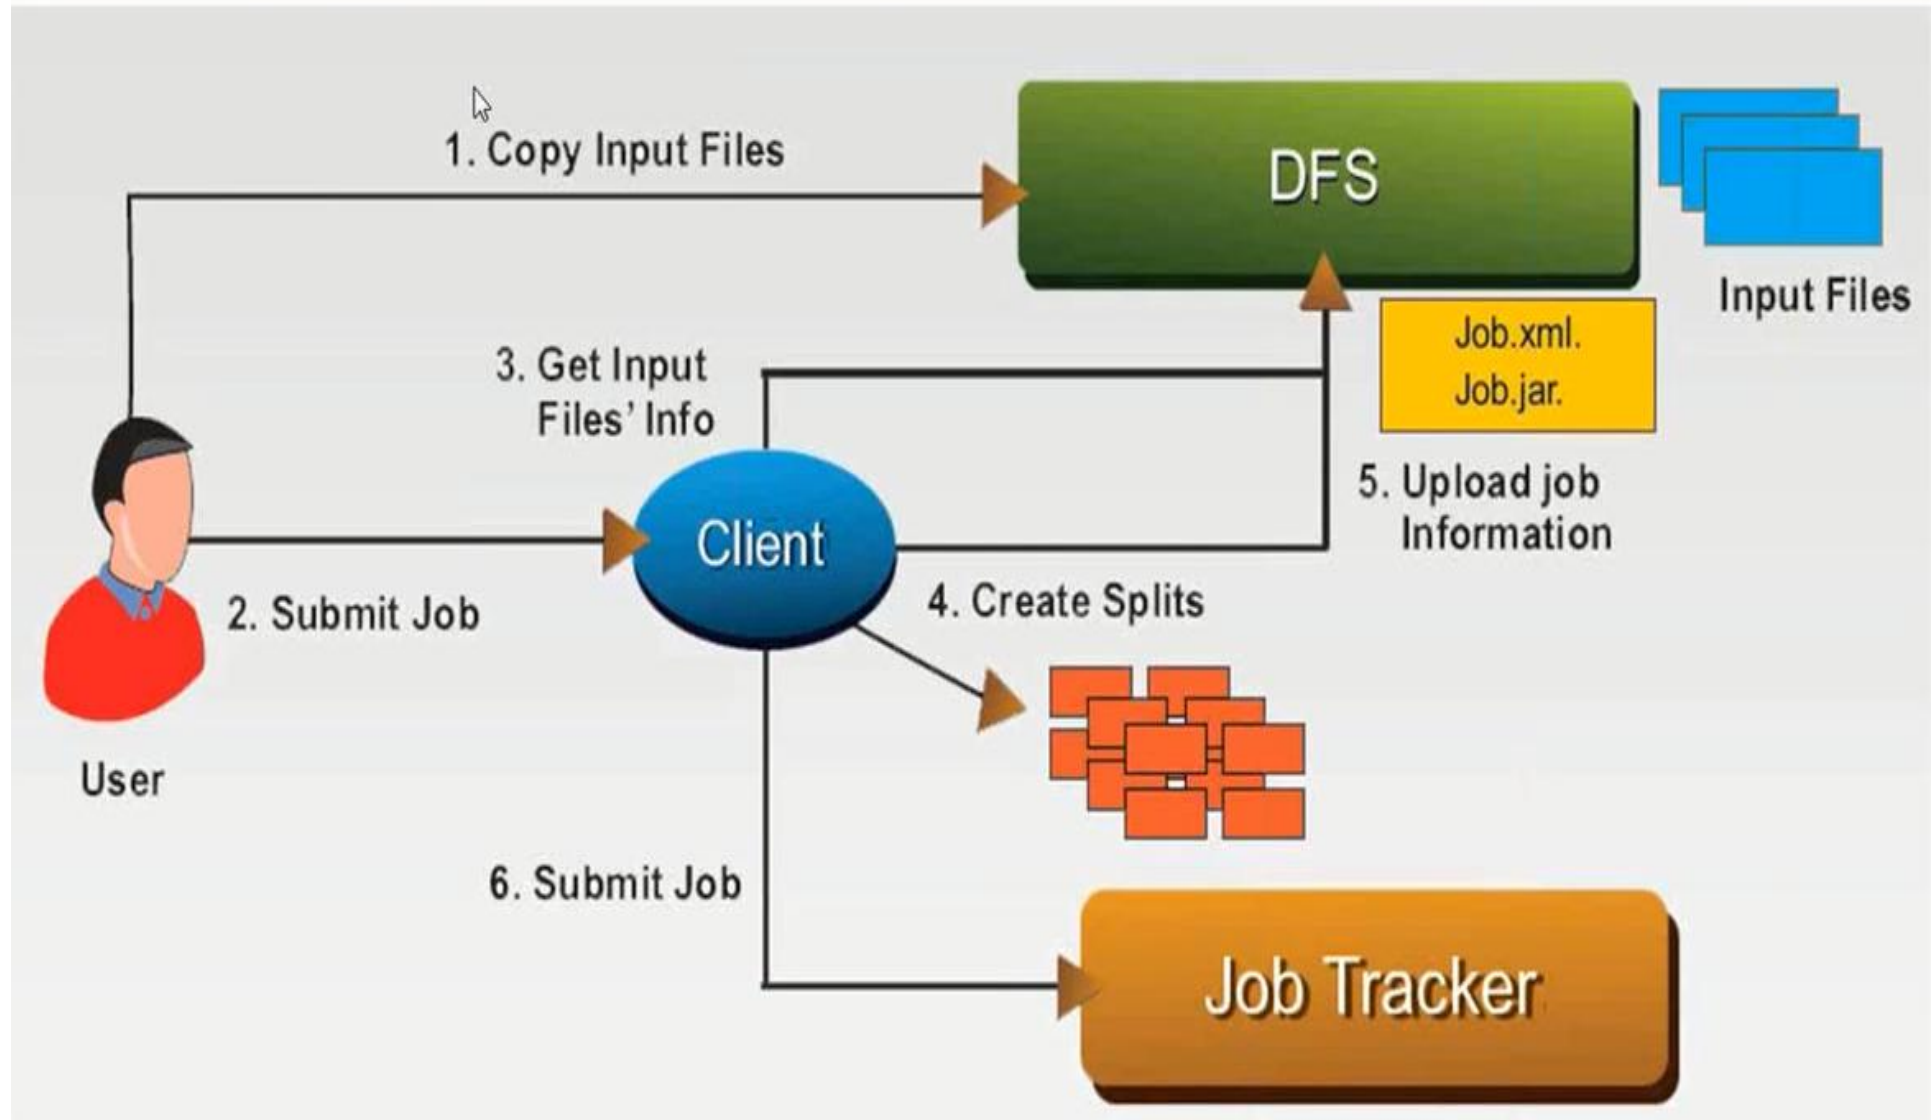
JOB TRACKER & TASK TRACKER



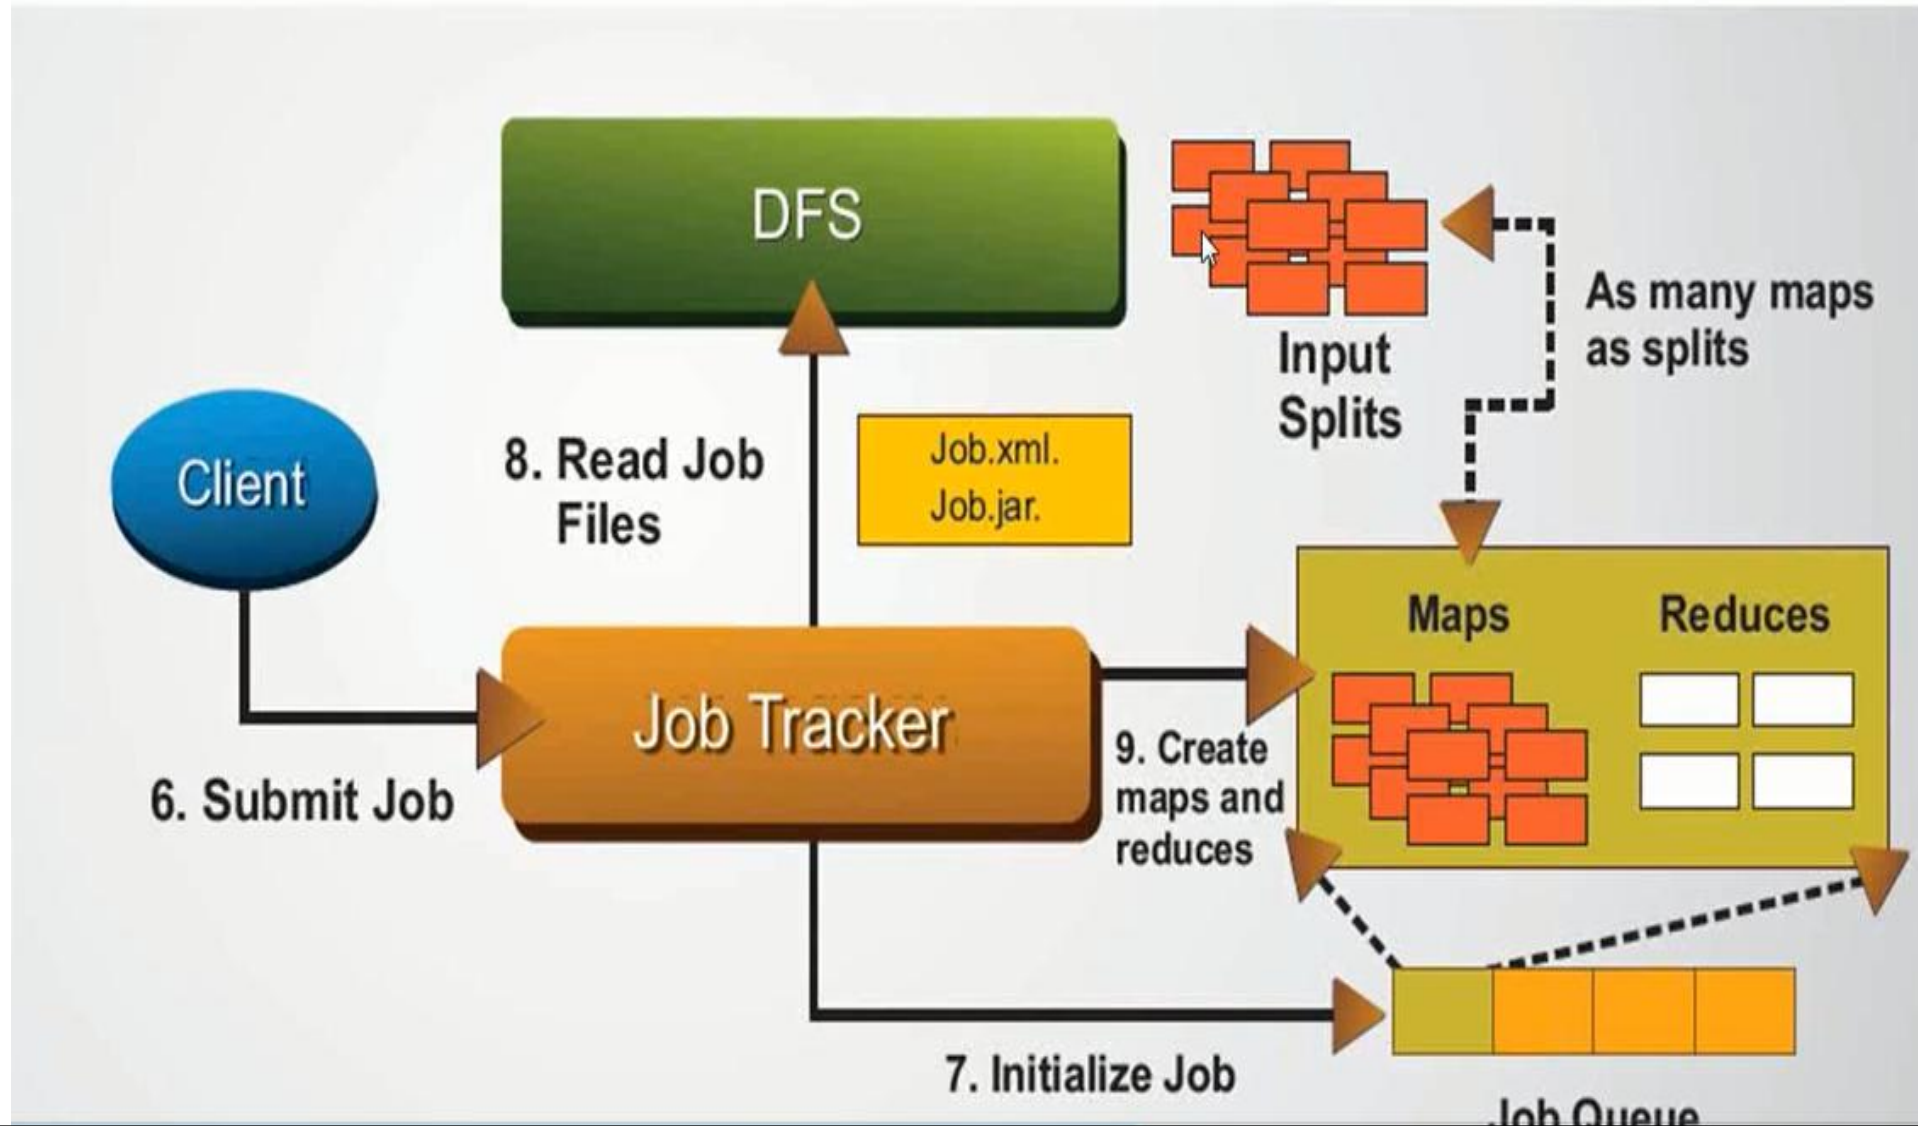
HDFS ARCHITECTURE



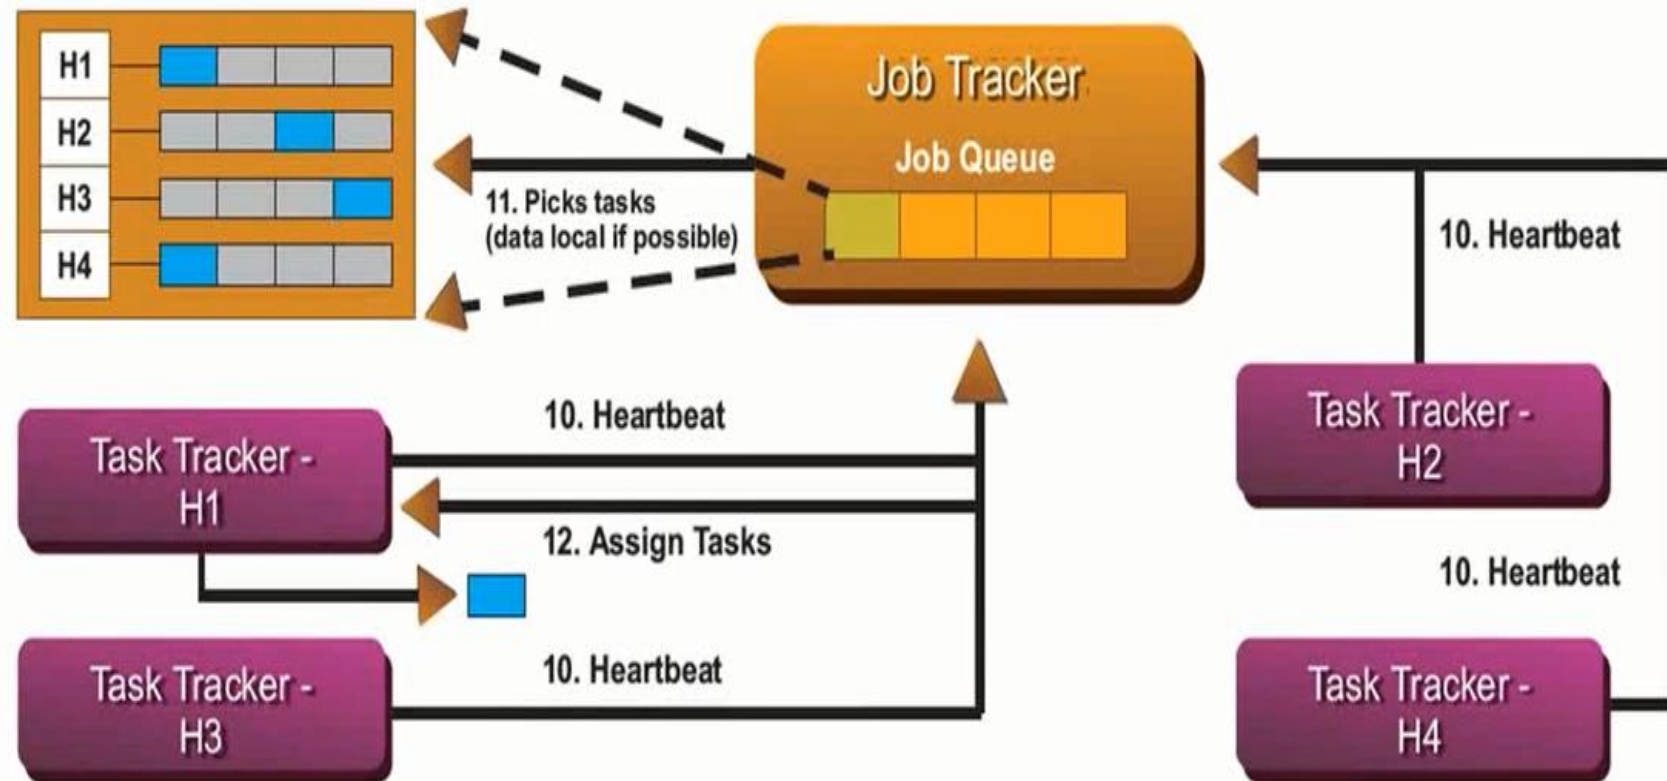
JOB TRACKER



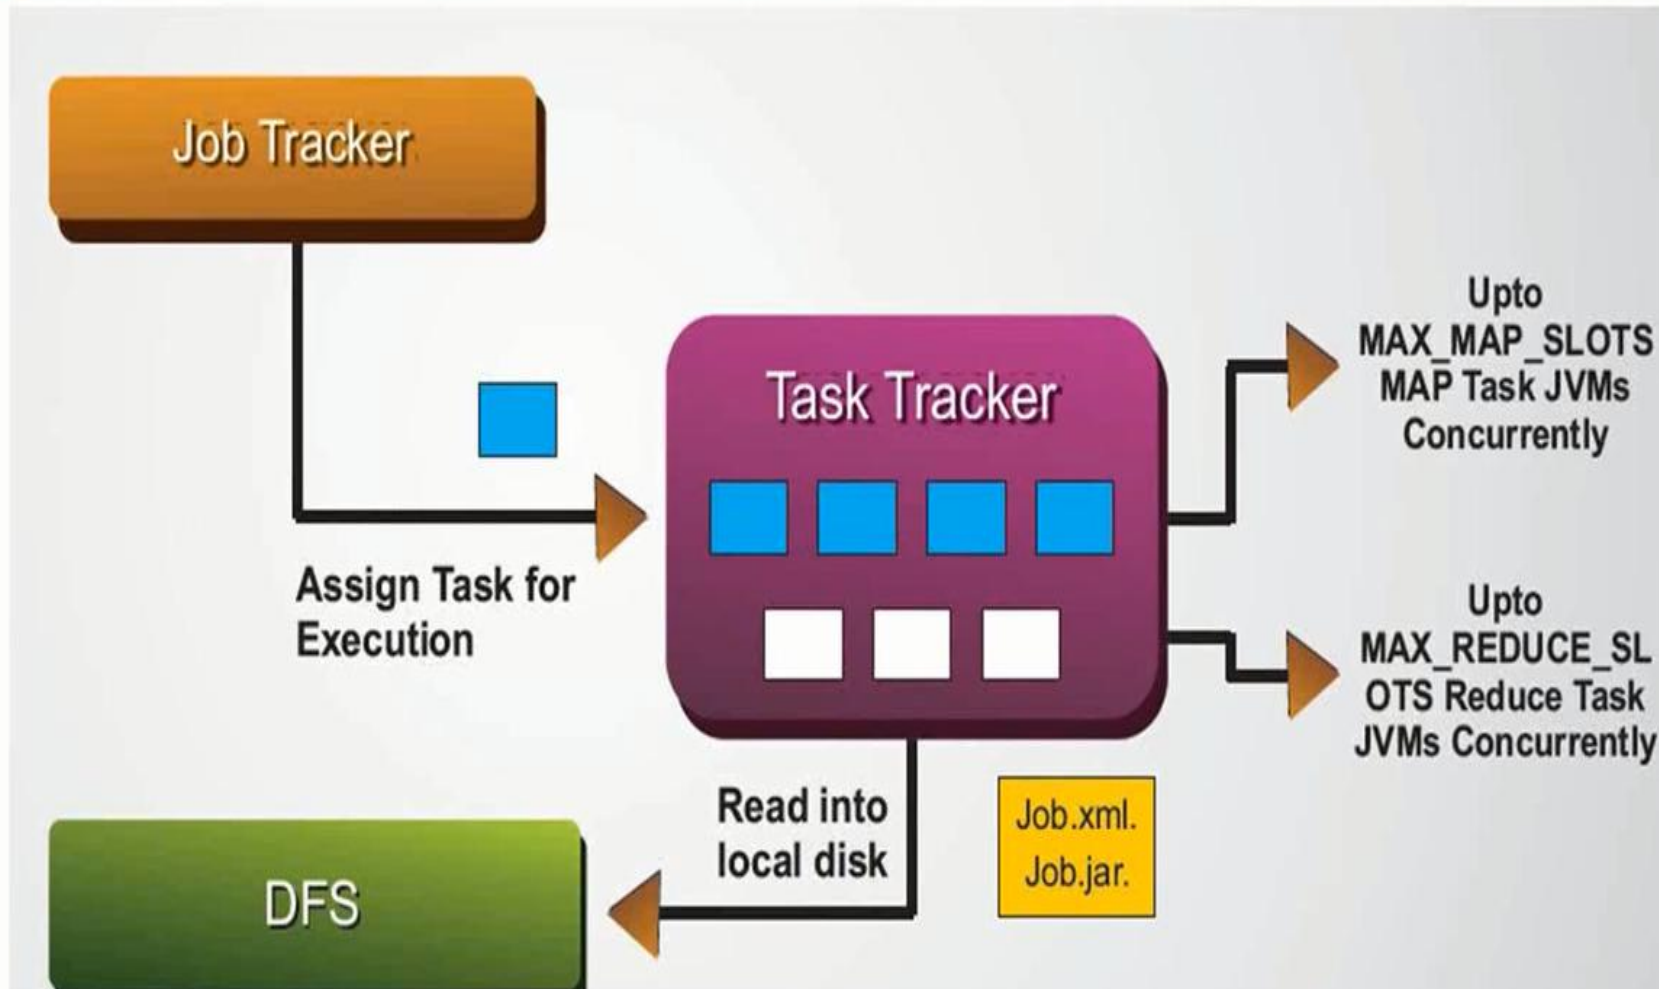
JOB TRACKER CONTD.....



JOB TRACKER CONTD....



JOB TRACKER CONTD.....



DATANODE

DataNode 1



DataNode 2



DataNode 3



DataNode 4



HADOOP DISTRIBUTED FILE SYSTEM

EXT4



BLOCK 1

BLOCK 3

EXT4



BLOCK 2

BLOCK 3

EXT4



BLOCK 3

BLOCK 1

BLOCK 2

EXT4



BLOCK 1

BLOCK 2

NAMENODE



aka Master

File Name - MyDatasetInHDFS

File Size - 350 MB

Replication Factor - 3

Blocks - BLK_0045732, BLK_9610590, BLK_8851209

Block Locations

Permissions

Created By, Created On

Last Modified By, Last Modified On

MyDatasetInHDFS

BLK_0045732	DN20	DN2	DN10
BLK_9610590	DN20	DN4	DN13
BLK_8851209	DN7	DN2	DN10


Name Node

In Disk - Metadata of Files & Folders

In Memory - Block locations



To All Data Nodes:
Send me block locations



HADOOP CLUSTER



NODE

CPU + RAM + DISK

Name Node Configuration

Processors: 2 Quad Core CPUs running @ 2 GHz

RAM: 128 GB

Disk: 6 x 1TB SATA

Network: 10 Gigabit Ethernet

Data Node Configuration

Processors: 2 Quad Core CPUs running @ 2 GHz

RAM: 64 GB

Disk: 12-24 x 1TB SATA

Network: 10 Gigabit Ethernet



Node




RACK

Rack
Collection of nodes



CLUSTER

Cluster
Racks interconnected over network





MAPREDUCE

SFO



SFO 5
SFO 3
SFO 4
SFO 1

SJOSE



SJSOE 2
SJSOE 6

LA



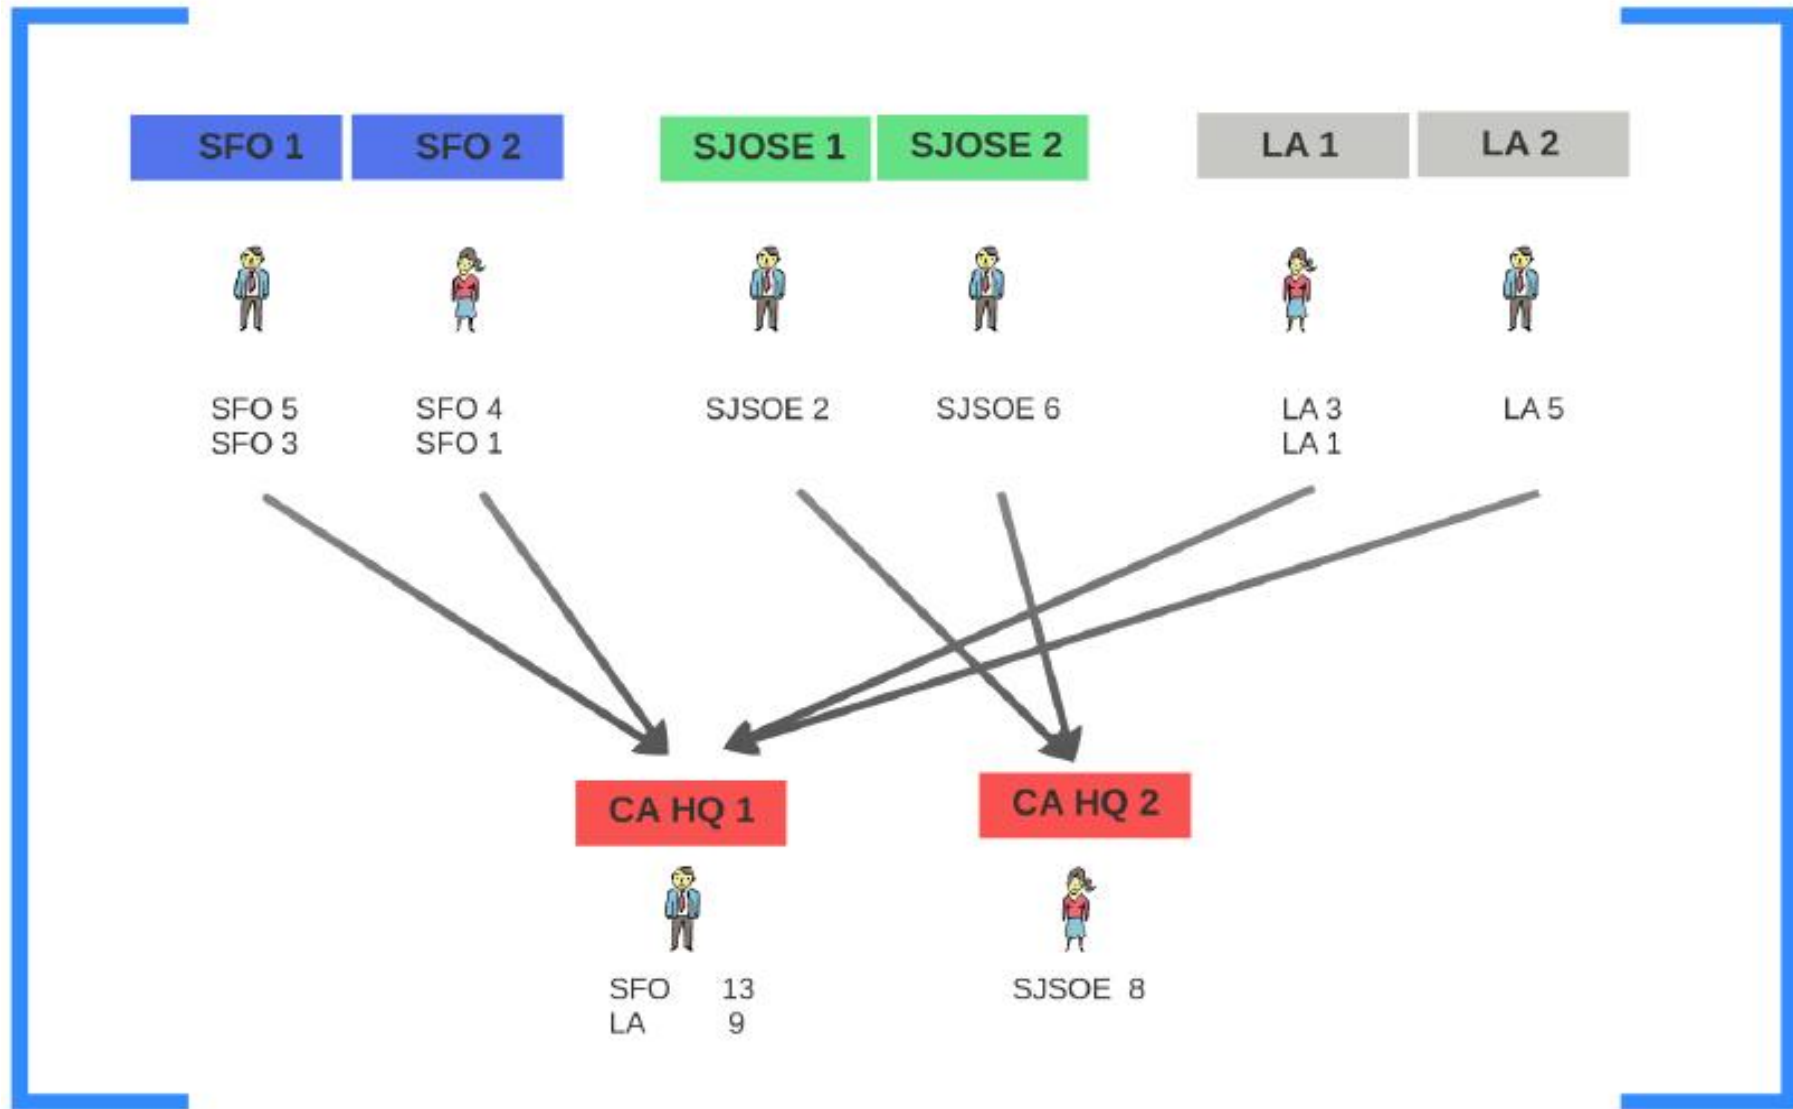
LA 3
LA 1
LA 5

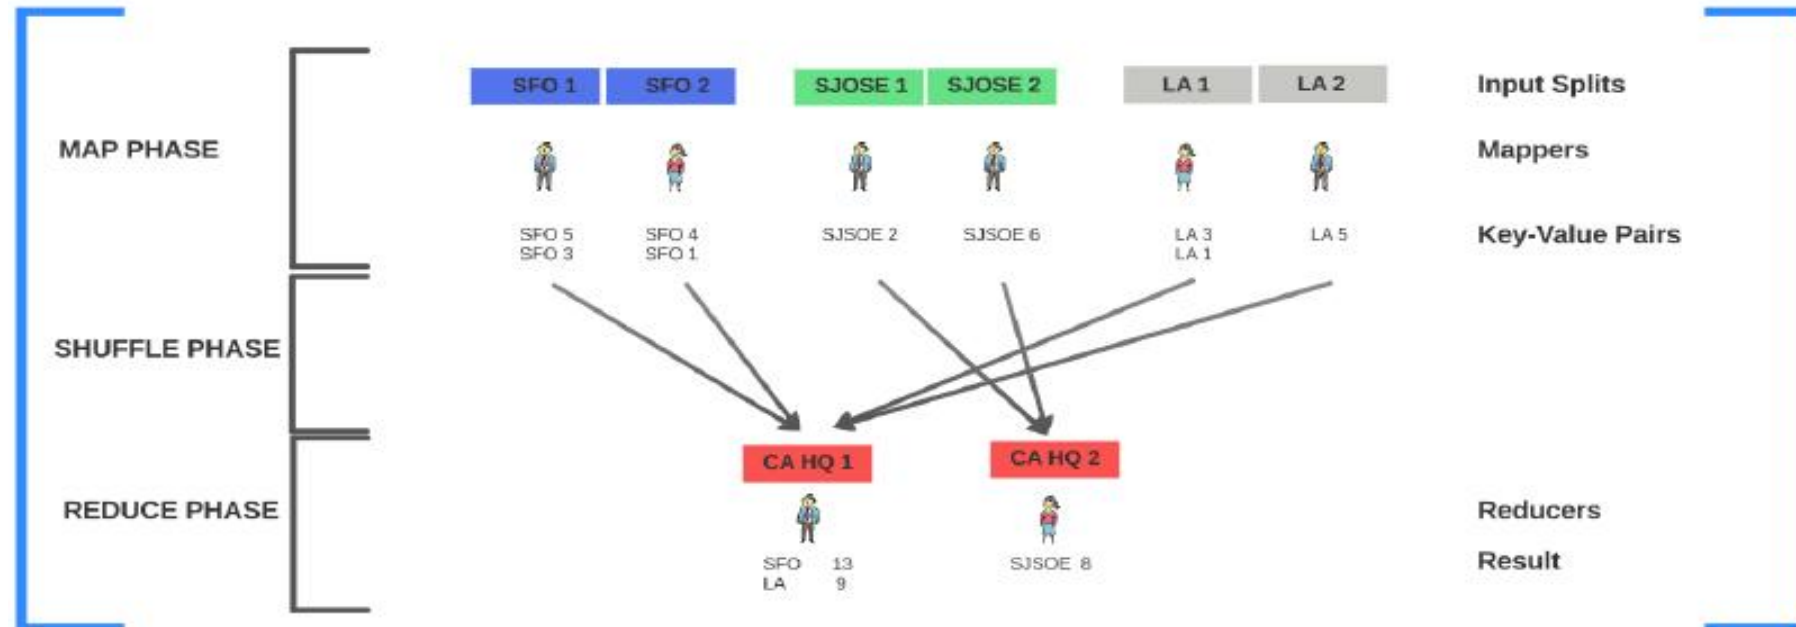


CA HQ



SFO	13
SJSOE	8
LA	9





WHAT IS MAPREDUCE?

- Distributed Programming model for processing large data sets
- Conceived at Google
- Can be implemented in any programming language
- MapReduce is NOT a programming language
- Hadoop implements MapReduce
- MapReduce System (Hadoop) - Manage communications, data transfers, parallel execution across distributed servers





DISSECTING MAPREDUCE COMPONENTS

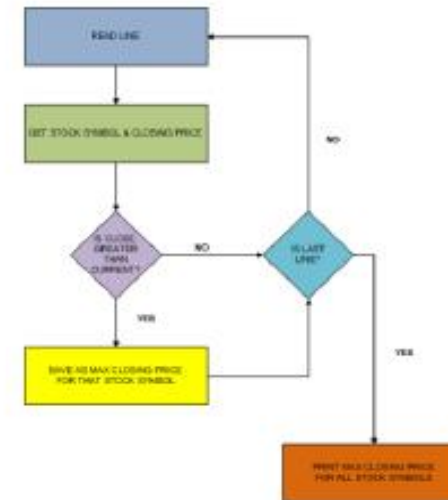
SAMPLE BIG DATA PROBLEM

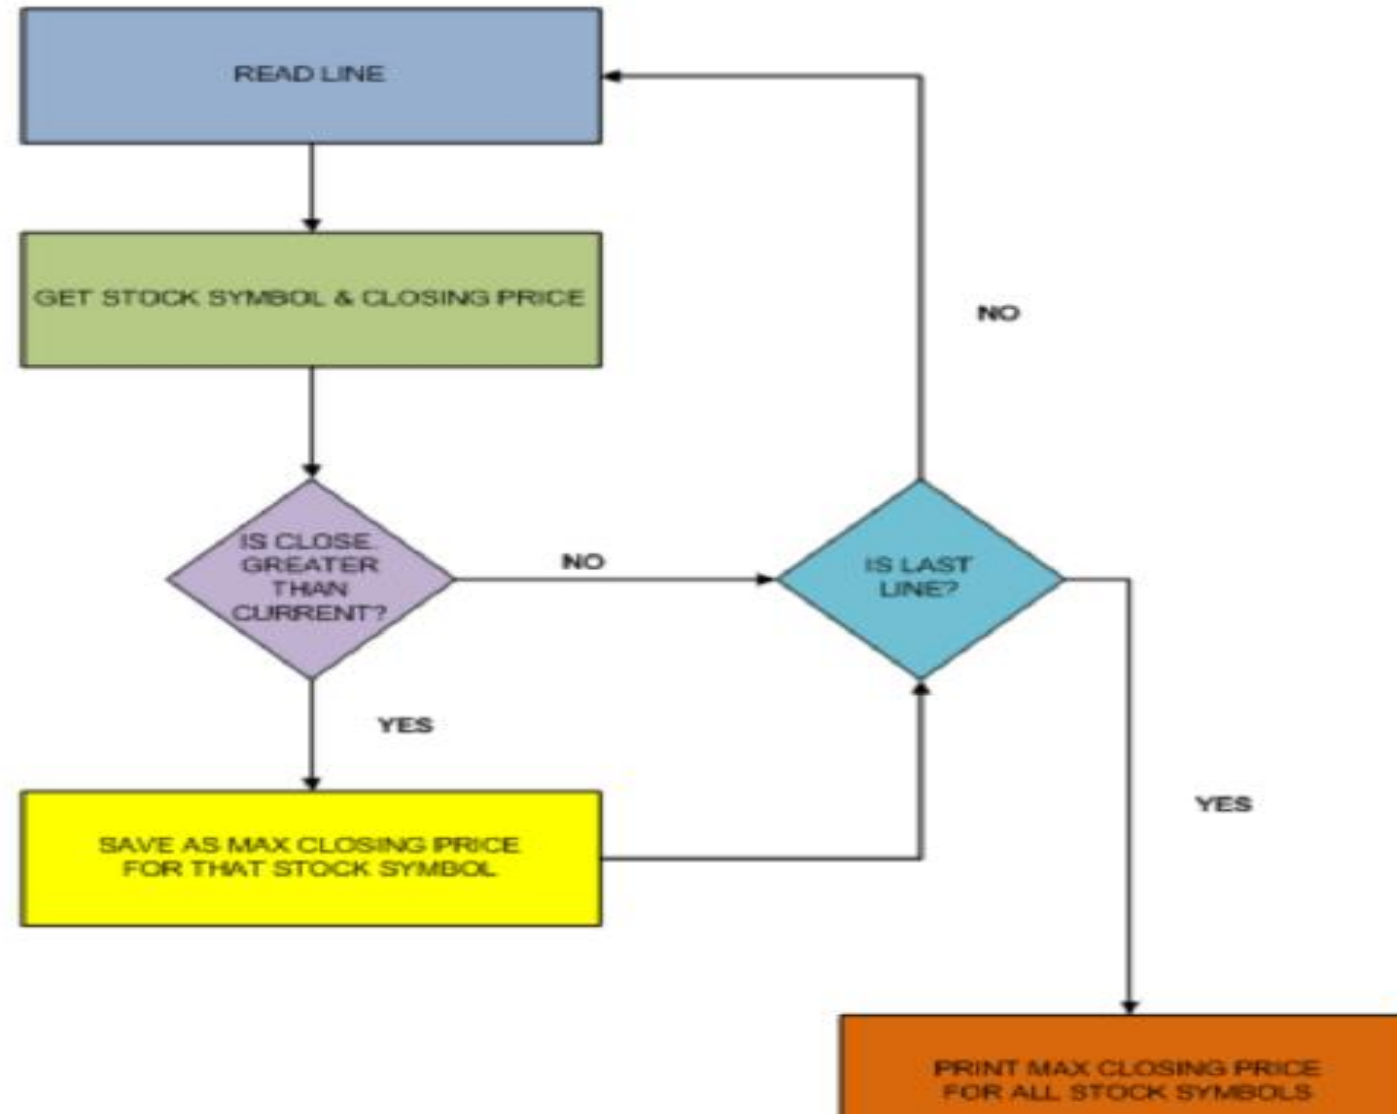
- Sample Stocks Dataset
- Each record has symbol, date, open, close...
- Find Maximum Closing Price for each symbol

```
ABCSE,B7J,2008-10-28,6.48,6.74,6.22,6.72,44300,5.79
ABCSE,B7J,2008-10-27,6.21,6.78,6.21,6.40,55200,5.51
ABCSE,B7J,2008-10-24,6.39,6.66,6.21,6.40,67400,5.51
ABCSE,B7J,2008-10-23,6.95,6.95,6.50,6.59,59400,5.68
ABCSE,B7J,2008-10-22,6.92,7.17,6.80,6.80,55300,5.86
ABCSE,B7J,2008-10-21,7.20,7.30,7.10,7.10,54400,6.11
ABCSE,B7J,2008-10-20,6.94,7.31,6.94,7.12,45700,6.13
ABCSE,B7J,2008-10-17,6.43,6.93,6.42,6.90,57700,5.94
ABCSE,B7J,2008-10-16,6.61,6.69,6.21,6.53,83200,5.62
ABCSE,B7J,2008-10-15,6.84,6.90,6.36,6.36,78900,5.48
ABCSE,B7J,2008-10-14,7.15,7.32,6.93,6.96,74700,5.99
ABCSE,B7J,2008-10-13,6.00,6.57,6.00,6.57,75700,5.66
ABCSE,B7J,2008-10-10,5.05,5.72,4.79,5.72,158400,4.93
ABCSE,B7J,2008-10-09,6.30,6.41,6.00,6.02,140500,5.18
ABCSE,B7J,2008-10-08,5.60,6.47,5.60,6.28,292000,5.41
ABCSE,B7J,2008-10-07,7.59,7.59,6.66,6.69,89900,5.76
ABCSE,B7J,2008-10-06,7.83,7.90,7.00,7.40,159600,6.37
```

MAX CLOSING PRICE ALGORITHM

- One Node
- Not Distributed

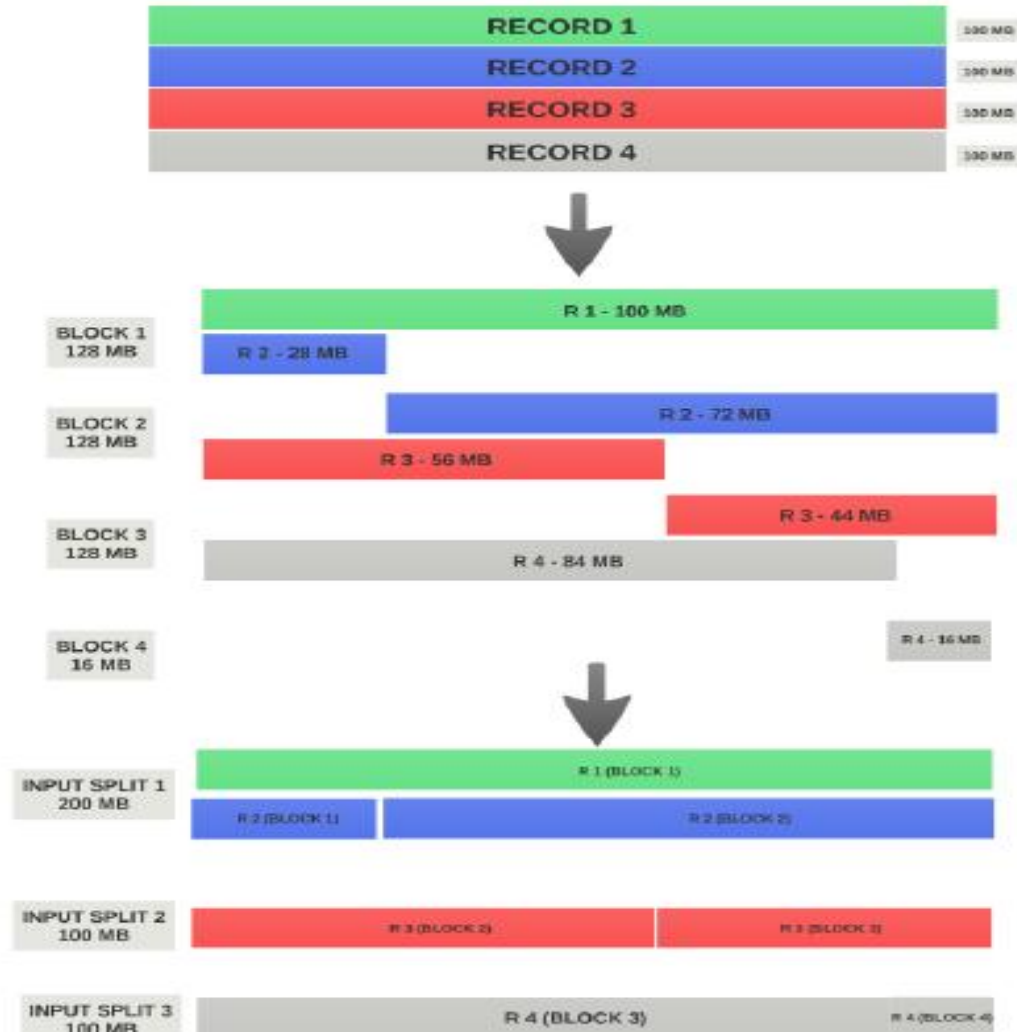




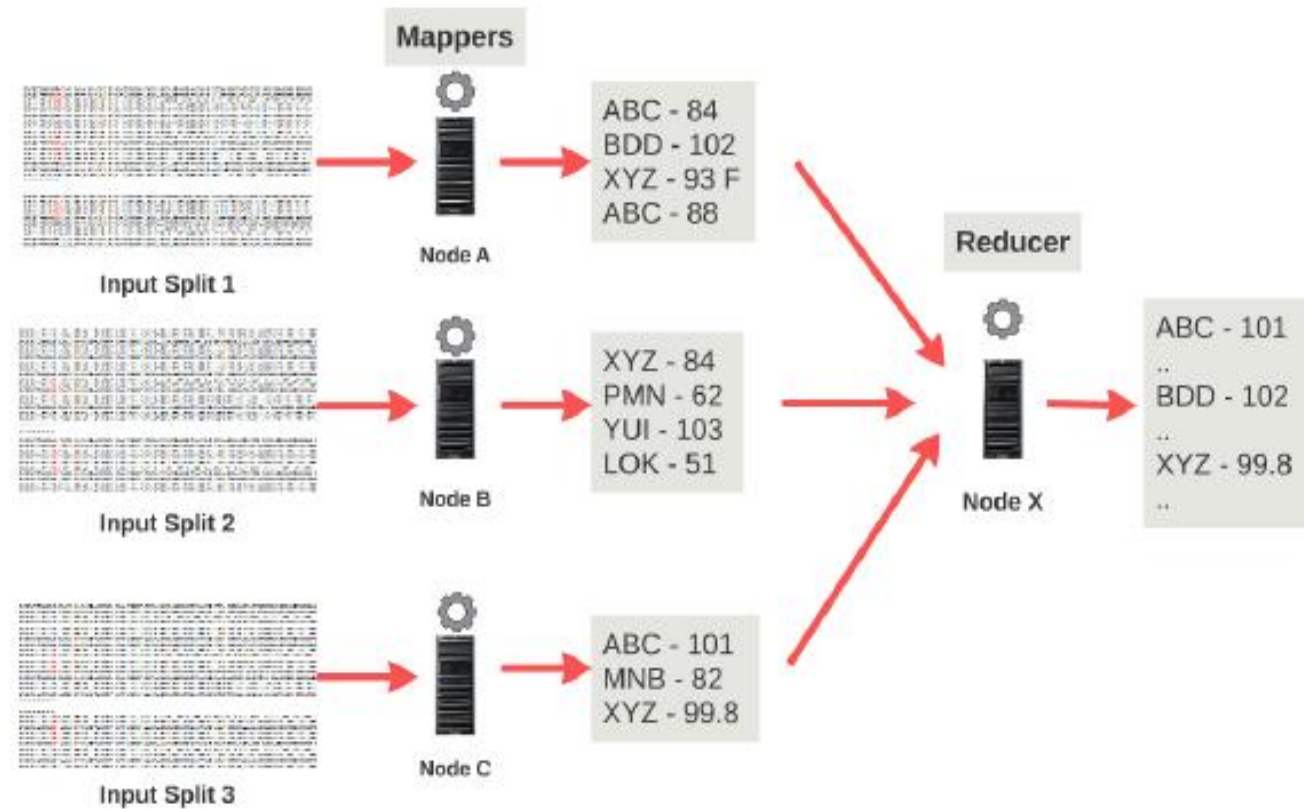
DISTRIBUTED



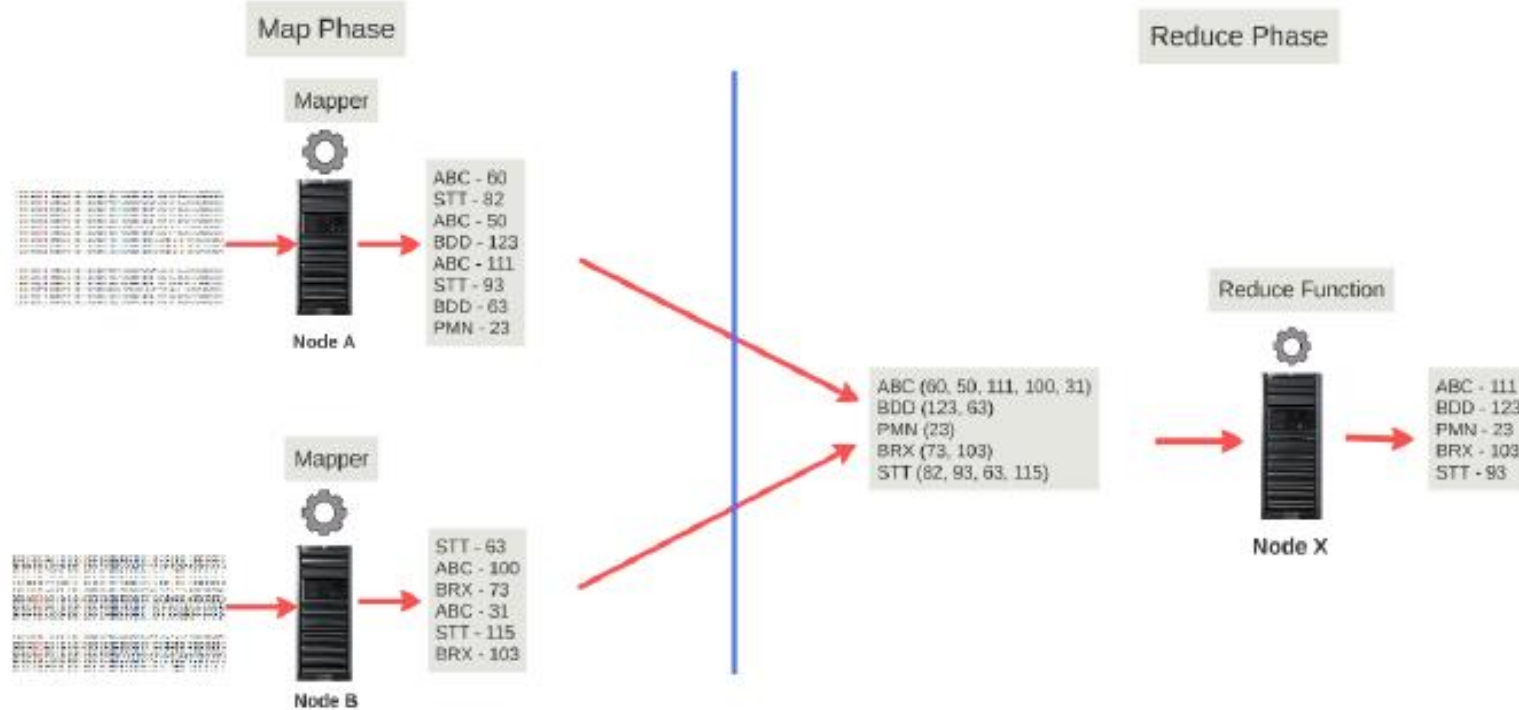
BLOCKS vs. INPUT SPLIT



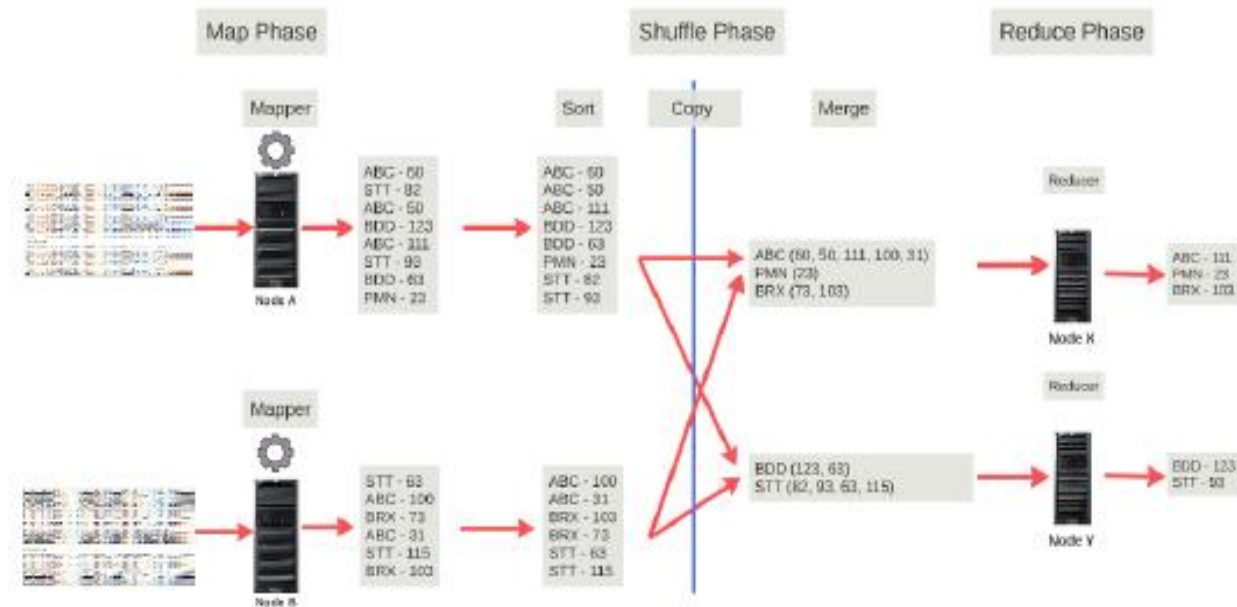
MAP PHASE



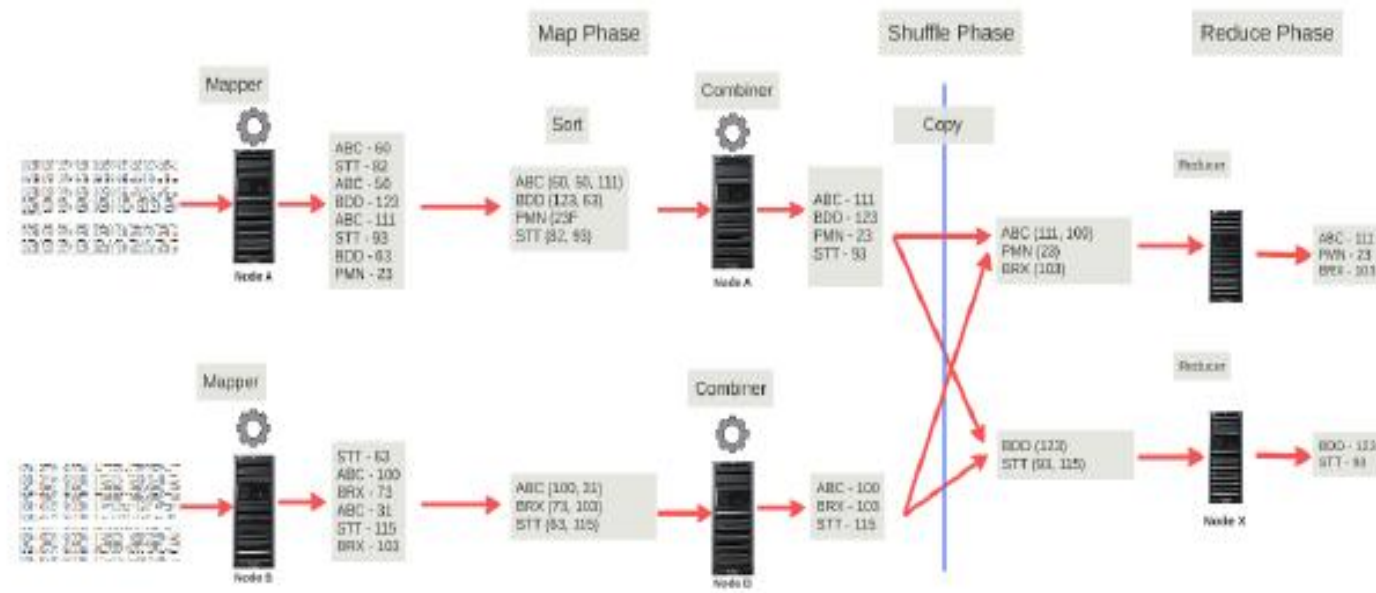
REDUCE PHASE



MULTIPLE REDUCERS



COMBINER (OPTIONAL)



STEPS FOR WRITING HADOOP MAP REDUCE PROGRAM

Step-1: Creating a File on Local System (file.txt)

```
$ cat > file.txt
```

Hi how are you

How is your class

How is your Big Data

What is time for lecture

What is the strength of Hadoop and Spark

Note: Use ctrl+d to save and make exit.

Step-2: Loading file.txt from local system to HDFS

```
$ hadoop-2.7.7/bin/hadoop fs -put file.txt file
```

Step-3: Writing Programs for Processing File

DriverCode.java

MapperCode.java

ReducerCode.java

Step-4: Compiling all .java Programs:

```
$javac -classpath $HADOOP_HOME/hadoop-core.jar *.java
```

Step-5: Creating jar Files:

```
$ jar cvf test.jar *.class
```

Step-6: Running test.jar on file (Already on HDFS)

```
$ hadoop jar test.jar DriverCode file TestHadoop
```

Word Count Program With Map Reduce and Java:

1. Download and Install Eclipse IDE

Eclipse is the most popular Integrated Development Environment (IDE) for developing Java applications. It is robust, feature-rich, easy-to-use and powerful IDE which is the #1 choice of almost Java programmers in the world. And it is totally FREE.

- As of now (fall 2018), the latest release of Eclipse is Photon. Click the following link to download Eclipse:

<http://www.eclipse.org/downloads/eclipse-packages>

- The following file will get downloaded for Linux:

`eclipse-committers-photon-R-linux-gtk.tar.gz`

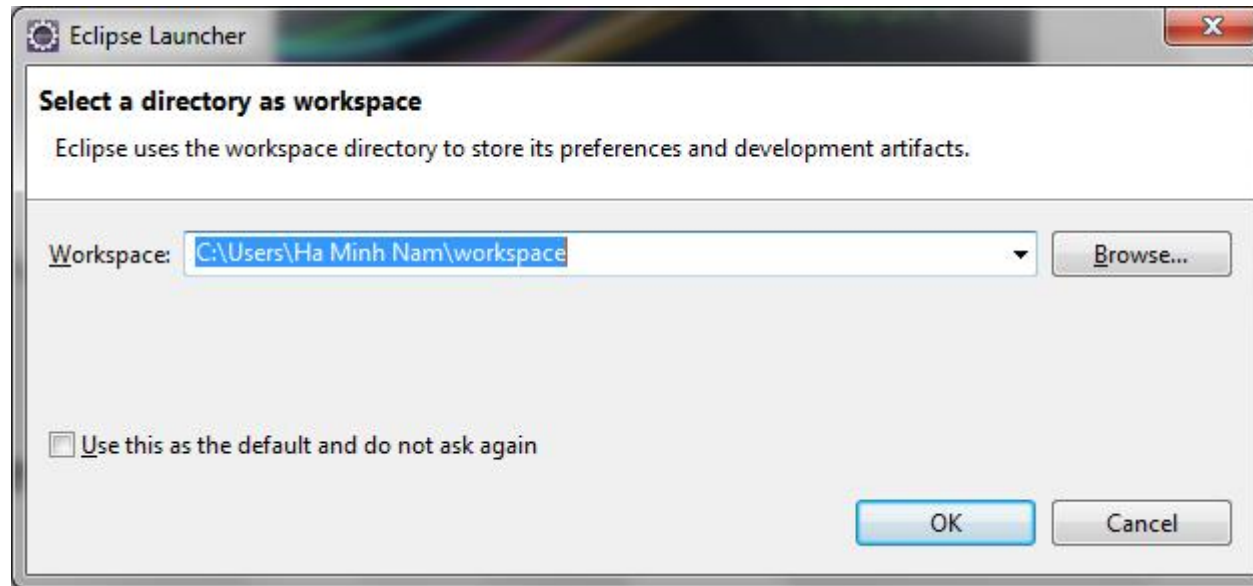
- Unzip the file using following command:

```
$ tar -zxvf eclipse-committers-photon-R-linux-gtk.tar.gz
```

- Move eclipse folder in Home Directory
- Open eclipse-> eclipse

2. Choose a Workspace Directory

Eclipse organizes projects by workspaces. A workspace is a group of related projects and it is actually a directory on your computer. That's why when you start Eclipse, it asks to choose a workspace location like this:



If you want to choose another directory, click **Browse**

Click **OK**. You should see the welcome screen:



3. Create Java Project in Package Explorer:

- File > New > Java Project >(Name it – WordCountProject) > Finish
- Right Click > New > Package (Name it - WordCountPackage) > Finish
- Right Click on Package > New > Class (Name it - WordCount)
- Add Following Reference Libraries –
 1. hadoop-core-1.2.1.jar
 2. commons-cli-1.2.jar
- Right Click on Project > Build Path> Add External Archivals
 - Downloads/hadoop-core-1.2.1.jar
 - Downloads/commons-cli-1.2.jar

4. Type following Java Program for Driver, Mapper and Reducer :

```
package WordCountPackage;  
import java.io.IOException;  
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;
```

```
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
```

```
public class WordCount {
    public static void main(String [] args) throws Exception
    {
        Configuration c=new Configuration();
        String[] files=new GenericOptionsParser(c,args).getRemainingArgs();
        Path input=new Path(files[0]);
        Path output=new Path(files[1]);
        Job j=new Job(c,"wordcount");
        j.setJarByClass(WordCount.class);
        j.setMapperClass(MapForWordCount.class);
        j.setReducerClass(ReduceForWordCount.class);
        j.setOutputKeyClass(Text.class);
        j.setOutputValueClass(IntWritable.class);
```

```
FileInputFormat.addInputPath(j, input);  
FileOutputFormat.setOutputPath(j, output);  
System.exit(j.waitForCompletion(true)?0:1);  
}
```

```
public static class MapForWordCount extends Mapper<LongWritable, Text, Text, IntWritable>{  
    public void map(LongWritable key, Text value, Context con) throws IOException, InterruptedException  
    {  
        String line = value.toString();  
        String[] words=line.split(" ");  
        for(String word: words )  
        {  
            Text outputKey = new Text(word.toUpperCase().trim());  
            IntWritable outputValue = new IntWritable(1);  
            con.write(outputKey, outputValue);  
        }  
    }  
}
```

```
public static class ReduceForWordCount extends Reducer<Text, IntWritable, Text, IntWritable>
{
    public void reduce(Text word, Iterable<IntWritable> values, Context con) throws IOException, InterruptedException
    {
        int sum = 0;
        for(IntWritable value : values)
        {
            sum += value.get();
        }
        con.write(word, new IntWritable(sum));
    }
}
```

The program consist of 3 classes:

- Driver class (Public void static main- the entry point)
- Map class which extends public class Mapper<KEYIN,VALUEIN,KEYOUT,VALUEOUT> and implements the Map function.
- Reduce class which extends public class Reducer<KEYIN,VALUEIN,KEYOUT,VALUEOUT> and implements the Reduce function.

5. Make Jar File:

Right Click on Project> Export> Java>Select export destination as **Jar File** > next> Finish

Jar File Path: /home/mamoon/temp.jar

6. Create a Text File and Move it to HDFS:

File Contents:

Apple, Banana, Mango,Apple,Mango, Banana, Apple, Banana, Mango,Apple,Mango, Banana

To Move this into Hadoop directly, open the terminal and enter the following commands:

```
$ hadoop-2.7.7/bin/hadoop fs -put /home/mamoon/fruits.txt demofruits.txt
```

7. Run Jar file

```
$ hadoop-2.7.7/bin/hadoop jar temp.jar WordCountPackage.WordCount demofruits.txt testfruitsoutput
```

8. Display Output:

```
$ hadoop-2.7.7/bin/hadoop fs -cat testfruitsoutput/part-r-00000
```

```
APPLE 7
```

```
BANANA 7
```

```
MANGO 7
```

Explanation of Map Reduce Computation:

Hadoop, MapReduce is a computation that decomposes large manipulation jobs into individual tasks that can be executed in parallel cross a cluster of servers. The results of tasks can be joined together to compute final results.

MapReduce consists of 2 steps:

Map Function – It takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (Key-Value pair).

Example – (Map function in Word Count)

<i>put</i>	<i>Set of data</i>	<i>Bus, Car, bus, car, train, car, bus, car, train, bus, TRAIN,BUS, buS, caR, CAR, car, BUS, TRAIN</i>
<i>Output</i>	<i>Convert into another set of data (Key,Value)</i>	<i>(Bus,1), (Car,1), (bus,1), (car,1), (train,1), (car,1), (bus,1), (car,1), (train,1), (bus,1), (TRAIN,1),(BUS,1), (buS,1), (caR,1), (CAR,1), (car,1), (BUS,1), (TRAIN,1)</i>

Reduce Function – Takes the output from Map as an input and combines those data tuples into a smaller set of tuples.
Example – (Reduce function in Word Count)

<i>Input (output of Map function)</i>	<i>Set of Tuples</i>	<i>(Bus,1), (Car,1), (bus,1), (car,1), (train,1), (car,1), (bus,1), (car,1), (train,1), (bus,1), (TRAIN,1),(BUS,1), (buS,1), (caR,1), (CAR,1), (car,1), (BUS,1), (TRAIN,1)</i>
<i>Output</i>	<i>Converts into smaller set of tuples</i>	<i>(BUS,7), (CAR,7), (TRAIN,4)</i>

Work Flow of Program

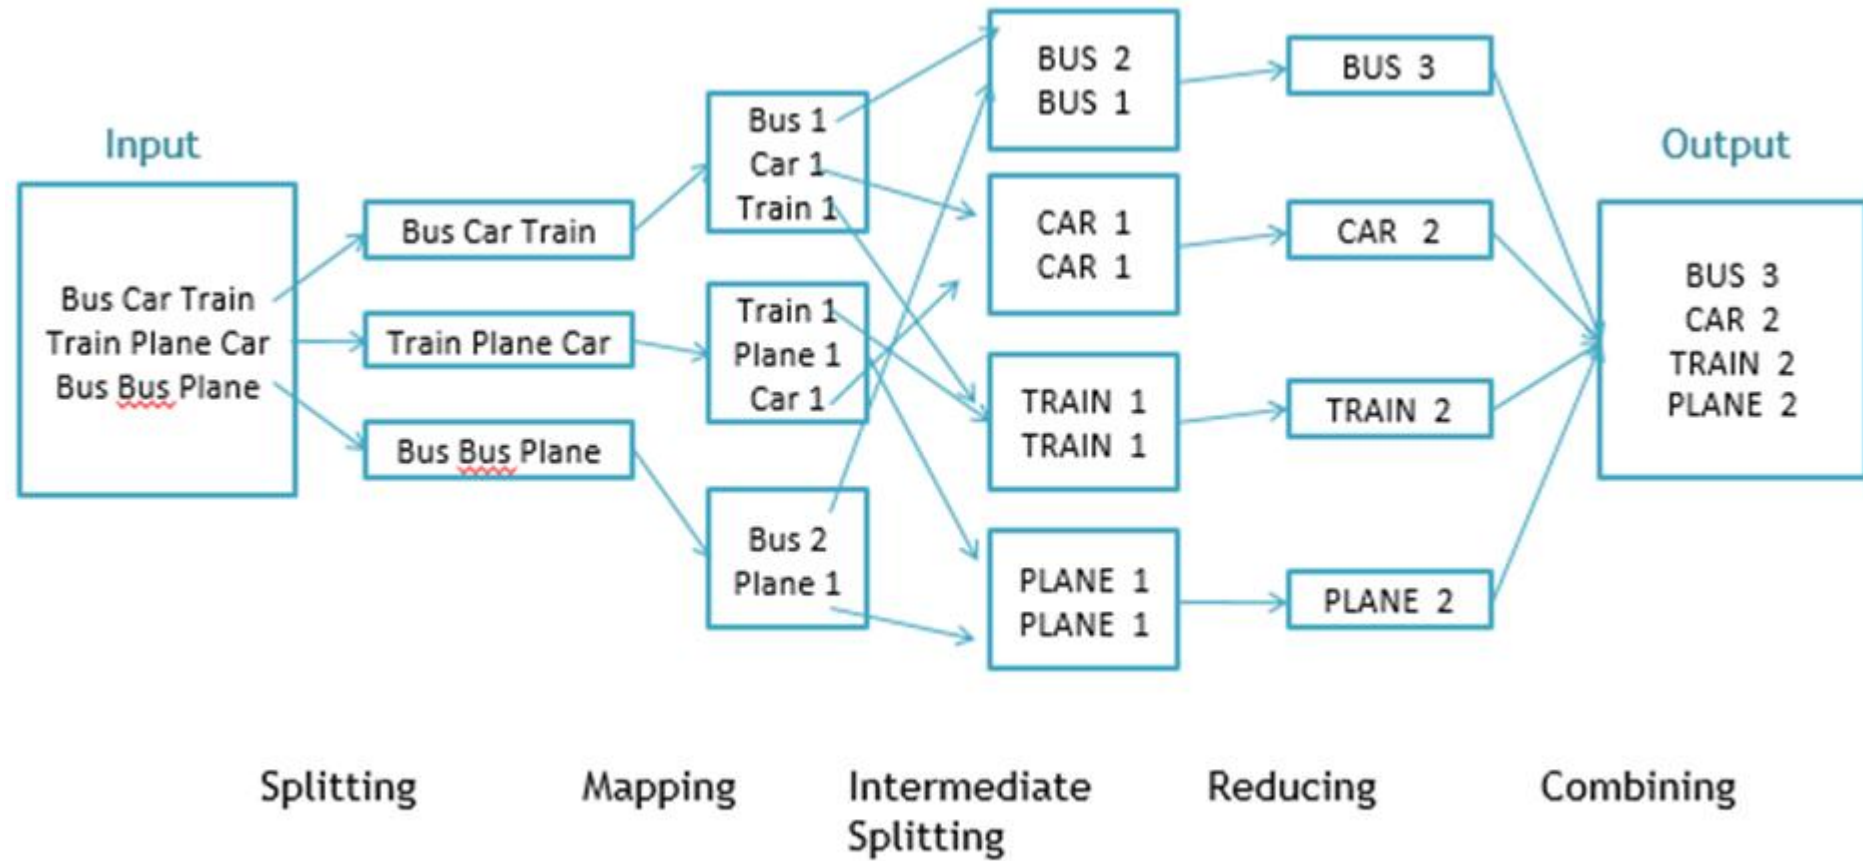


Fig. WorkFlow of MapReducing

Workflow of MapReduce consists of 5 steps:

Splitting – The splitting parameter can be anything, e.g. splitting by space, comma, semicolon, or even by a new line ('\n').

Mapping – as explained above

Intermediate splitting – the entire process in parallel on different clusters. In order to group them in “Reduce Phase” the similar KEY data should be on same cluster.

Reduce – it is nothing but mostly group by phase

Combining – The last phase where all the data (individual result set from each cluster) is combine together to form a Result

MapReduce Design Pattern:

To solve any problem in MapReduce, we need to think in terms of MapReduce. It is not necessarily true that every time we have both a map and reduce job.

Input-Map-Reduce-Output

Input-Map-Output

Input-Multiple Maps-Reduce-Output

Input-Map-Combiner-Reduce-Output

Following is a real time scenario to understand when to use which design pattern.


- Input-Map-Reduce-Output



If we want to do some aggregation then this pattern is used:

<i>Scenario</i>	<i>Counting gender total/ average salary of employees</i>
<i>Map (Key, Value)</i>	<i>Key: Gender Value: Their Salary</i>
<i>Reduce</i>	<i>Group by Gender And Take Total of salary for each group</i>


- **Input-Map-Output:**



If we want to change only the format of data then this pattern is used:

Scenario	Some employees have gender entry as "Female", "F", "f", "O"
Map (Key, Value)	Key : Employee Id Value : Gender -> if Gender is Female/ F/ f/ O then converted to F else if Gender is Male/M/m/1 then convert to M

• **Input-Multiple Maps-Reduce-Output**



In this design pattern, our input is taken from two files which have a different schema:

We have to find the total of gender-wide salary. But he have 2 files with different schema.

Input File 1

Gender is given as a prefix to Name

Eg. Ms. Shital Katkar

Mr. Krishna Katkar

Input File 2

There is a different column for gender. However, the format is mixed.

Eg. Female/Male, 0/1, F/M

Map (Key, Value)

Map 1 (For input 1)

We need to write a program to split prefix from Name and, according to the prefix, determine the gender.

Then prepare the Key value pair (Gender,Salary).

Map 2 (For input 2)


Here the program will be straightforward. Resolve the mixed format and make a key value pair (Gender,Salary).

Reduce

Group by Gender

And Take Total of salary for each group

- Input-Map-Combiner-Reduce-Output:



A Combiner, also known as a semi-reducer, is an optional class that operates by accepting the inputs from the Map class and thereafter passing the output key-value pairs to the Reducer class. Purpose of the combiner is to reduce workload of Reducer.

In MapReduce program, 20% of the work is done in the Map Stage, which is also known as the data preparation stage, which works in parallel.

80% of the work is done in Reduce stage which is known as the calculation stage, and it is not parallel. Therefore it is slower than the Map phase. To reduce time, some work in the Reduce phase can be done in the combiner phase.

Scenario

There are 5 departments. And we have to calculate the gender-wide total salary. However, there are certain rules to calculate the total. After calculating gender wise total for each department, if the salary is greater than 200K, add 20K in total, if the salary is greater than 100K, add 10K in total

Input Files (for each department there is 1 file)	Map (Parallel) (, Value = Salary)	Combiner (Parallel)	Reducer (Not Parallel)	Output
Dept 1	Male<10,20,25,45,15,45,25,20> Female <10,30,20,25,35>	Male <250,20> Female <120,10>	Male < 250,20,155, 10,90,90,30> Female <120,10,175,10,135, 10,110,10,130,10>	Male <645> Female <720>
Dept 2	Male<15,30,40,25,45> Female <20,35,25,35,40>	Male <155,10> Female <175,10>		
Dept 3	Male<10,20,20,40> Female <10,30,25,70>	Male <90,00> Female <135,10>		
Dept 4	Male<45,25,20> Female <30,20,25,35>	Male <90,00> Female <110,10>		
Dept 5	Male<10,20> Female <10,30,20,25,35>	Male <30,00> Female <130,10>		